

Capstone: Predict Adult Income

Introduction

The goal of this project is to predict whether the income of adults is greater than \$50k/yr. This is a classification problem. We use a couple of algorithms to train the model and then choose the most suitable one.

Dataset

Adult Data Set is used for the prediction purpose. The dataset can be downloaded using this URL: [\[https://archive.ics.uci.edu/ml/datasets/Adult\]](https://archive.ics.uci.edu/ml/datasets/Adult). As mentioned in this website - “Extraction was done by Barry Becker from the 1994 Census database. A set of reasonably clean records was extracted using the following conditions: ((AAGE>16) && (AGI>100) && (AFNLWGT>1)&& (HRSWK>0)). Prediction task is to determine whether a person makes over 50K a year.”

The following files are available:

1. adult.data
2. adult.test
3. adult.names

adult.names file provides us with information about the data. It is mentioned that there are a total of 48842 instances which is split into train-test datasets. Train dataset is available in adult.data file, it has 32561 instances and test dataset is available in adult.test file, it has 16281 instances.

The variables of this dataset are as follows:

age, workclass, fnlwgt, education, education_num, marital_status, occupation, relationship, race, sex, capital_gain, capital_loss, hours_per_week, native_country, income - variable to be predicted.

Procedure

First we download data and then preprocess it. Data exploration is done to identify which of the above mentioned variables are to be used for income prediction. Then the following algorithms are applied to train the model. For training the model, data is split into train(90%) and validation(10%) datasets.

Algorithms used -

1. Logistic Regression
2. Support Vector Classifier
3. Random Forest Classifier
4. Gradient Boosting Classifier

Model performance is evaluated based on Overall Accuracy and F1 score. Confusion Matrix is used to obtain Accuracy and F1 score.

Accuracy - This is defined as the no. of correct predictions divided by the total number of the dataset.

F1-score - This is the harmonic mean of precision and recall.

$$F1 - score = \frac{2 * precision * recall}{precision + recall}$$

where, precision is Positive Predictive Value recall is sensitivity

If results of confusion Matrix is stored in “cm”, we can obtain accuracy using the command `cm$overall[“Accuracy”]` and F1-score using the command `cm$byClass[“F1”]`

Downloading Data:

The required libraries and data are downloaded. Data is read into the file “income_data” and column names are assigned to the data.

```
#-----  
#Download required packages  
#-----  
  
if(!require(tidyverse))  
  install.packages("tidyverse", repos = "http://cran.us.r-project.org")  
if(!require(caret))  
  install.packages("caret", repos = "http://cran.us.r-project.org")  
if(!require(data.table))  
  install.packages("data.table", repos = "http://cran.us.r-project.org")  
if(!require(e1071))  
  install.packages("e1071", repos = "http://cran.us.r-project.org")  
if(!require(randomForest))  
  install.packages("randomForest", repos = "http://cran.us.r-project.org")  
  
#-----  
# Part 1: Downloading data, preprocessing and Data exploration  
#-----  
  
# Downloading and Reading data  
  
if(!file.exists("adult.data")){  
  download.file("http://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data", "adult.data")  
}  
income_data <- read.csv("adult.data", header = FALSE)  
  
# Set column names  
colnames(income_data) <- c("age", "workclass", "fnlwgt", "education", "education_num", "marital_status",  
                           "occupation", "relationship", "race", "sex", "capital_gain", "capital_loss",  
                           "hours_per_week", "native_country", "income")
```

Methods/Analysis

Data preprocessing

First we check the data dimensions. There are 32561 rows and 15 columns. The structure of data reveals that there are 6 columns of type integer and 9 columns of type factor. Also, we can notice some columns have data in the form " ?" which is unknown/missing data(mentioned in adult.names) file. We can further confirm this using summary() and levels command. The no. of rows with missing data is few in number, so we remove these rows.

We re-read the income_data file and convert " ?" to NA and then omit rows containing NA. We set the column names again. Now the no. of rows is 30162 and 15 columns.

```
# Explore data  
  
# To get the dimension of data  
dim(income_data)
```

```
## [1] 32561    15
```

```
# To check the structure of data
str(income_data)
```

```
## 'data.frame': 32561 obs. of 15 variables:
## $ age : int 39 50 38 53 28 37 49 52 31 42 ...
## $ workclass : Factor w/ 9 levels " ?"," Federal-gov",...: 8 7 5 5 5 5 7 5 5 ...
## $ fnlwt : int 77516 83311 215646 234721 338409 284582 160187 209642 45781 159449 ...
## $ education : Factor w/ 16 levels " 10th"," 11th",...: 10 10 12 2 10 13 7 12 13 10 ...
## $ education_num : int 13 13 9 7 13 14 5 9 14 13 ...
## $ marital_status: Factor w/ 7 levels " Divorced"," Married-AF-spouse",...: 5 3 1 3 3 3 4 3 5 3 ...
## $ occupation : Factor w/ 15 levels " ?"," Adm-clerical",...: 2 5 7 7 11 5 9 5 11 5 ...
## $ relationship : Factor w/ 6 levels " Husband"," Not-in-family",...: 2 1 2 1 6 6 2 1 2 1 ...
## $ race : Factor w/ 5 levels " Amer-Indian-Eskimo",...: 5 5 5 3 3 5 3 5 5 5 ...
## $ sex : Factor w/ 2 levels " Female"," Male": 2 2 2 2 1 1 1 2 1 2 ...
## $ capital_gain : int 2174 0 0 0 0 0 0 0 14084 5178 ...
## $ capital_loss : int 0 0 0 0 0 0 0 0 0 0 ...
## $ hours_per_week: int 40 13 40 40 40 40 16 45 50 40 ...
## $ native_country: Factor w/ 42 levels " ?"," Cambodia",...: 40 40 40 40 6 40 24 40 40 40 ...
## $ income : Factor w/ 2 levels " <=50K"," >50K": 1 1 1 1 1 1 1 2 2 2 ...
```

```
# To get summary of data
summary(income_data)
```

```
##      age      workclass      fnlwt
## Min.   :17.00   Private      :22696   Min.    : 12285
## 1st Qu.:28.00   Self-emp-not-inc: 2541   1st Qu.: 117827
## Median :37.00   Local-gov       : 2093   Median : 178356
## Mean   :38.58   ?               : 1836   Mean   : 189778
## 3rd Qu.:48.00   State-gov       : 1298   3rd Qu.: 237051
## Max.   :90.00   Self-emp-inc    : 1116   Max.    :1484705
##              (Other)      : 981
##      education      education_num      marital_status
## HS-grad      :10501   Min.    : 1.00   Divorced      : 4443
## Some-college: 7291   1st Qu.: 9.00   Married-AF-spouse : 23
## Bachelors    : 5355   Median :10.00   Married-civ-spouse :14976
## Masters      : 1723   Mean    :10.08   Married-spouse-absent: 418
## Assoc-voc    : 1382   3rd Qu.:12.00   Never-married    :10683
## 11th         : 1175   Max.    :16.00   Separated        : 1025
## (Other)      : 5134           Widowed          : 993
##      occupation      relationship
## Prof-specialty :4140   Husband         :13193
## Craft-repair   :4099   Not-in-family   : 8305
## Exec-managerial:4066   Other-relative: 981
## Adm-clerical   :3770   Own-child       : 5068
## Sales          :3650   Unmarried       : 3446
## Other-service  :3295   Wife            : 1568
## (Other)        :9541
##      race      sex      capital_gain
## Amer-Indian-Eskimo: 311   Female:10771   Min.    : 0
## Asian-Pac-Islander: 1039   Male  :21790   1st Qu.: 0
## Black              : 3124           Median : 0
## Other              : 271           Mean   : 1078
## White              :27816           3rd Qu.: 0
##                      Max.    :99999
##
```

```
## capital_loss hours_per_week native_country income
## Min. : 0.0 Min. : 1.00 United-States:29170 <=50K:24720
## 1st Qu.: 0.0 1st Qu.:40.00 Mexico : 643 >50K : 7841
## Median : 0.0 Median :40.00 ? : 583
## Mean : 87.3 Mean :40.44 Philippines : 198
## 3rd Qu.: 0.0 3rd Qu.:45.00 Germany : 137
## Max. :4356.0 Max. :99.00 Canada : 121
## (Other) : 1709
```

```
# To explore the levels in each column
sapply(income_data, levels)
```

```
## $age
## NULL
##
## $workclass
## [1] " ?" " Federal-gov" " Local-gov"
## [4] " Never-worked" " Private" " Self-emp-inc"
## [7] " Self-emp-not-inc" " State-gov" " Without-pay"
##
## $fnlwgt
## NULL
##
## $education
## [1] " 10th" " 11th" " 12th" " 1st-4th"
## [5] " 5th-6th" " 7th-8th" " 9th" " Assoc-acdm"
## [9] " Assoc-voc" " Bachelors" " Doctorate" " HS-grad"
## [13] " Masters" " Preschool" " Prof-school" " Some-college"
##
## $education_num
## NULL
##
## $marital_status
## [1] " Divorced" " Married-AF-spouse"
## [3] " Married-civ-spouse" " Married-spouse-absent"
## [5] " Never-married" " Separated"
## [7] " Widowed"
##
## $occupation
## [1] " ?" " Adm-clerical" " Armed-Forces"
## [4] " Craft-repair" " Exec-managerial" " Farming-fishing"
## [7] " Handlers-cleaners" " Machine-op-inspct" " Other-service"
## [10] " Priv-house-serv" " Prof-specialty" " Protective-serv"
## [13] " Sales" " Tech-support" " Transport-moving"
##
## $relationship
## [1] " Husband" " Not-in-family" " Other-relative" " Own-child"
## [5] " Unmarried" " Wife"
##
## $race
## [1] " Amer-Indian-Eskimo" " Asian-Pac-Islander" " Black"
## [4] " Other" " White"
##
## $sex
## [1] " Female" " Male"
```

```
##
## $capital_gain
## NULL
##
## $capital_loss
## NULL
##
## $hours_per_week
## NULL
##
## $native_country
## [1] " ?" " Cambodia"
## [3] " Canada" " China"
## [5] " Columbia" " Cuba"
## [7] " Dominican-Republic" " Ecuador"
## [9] " El-Salvador" " England"
## [11] " France" " Germany"
## [13] " Greece" " Guatemala"
## [15] " Haiti" " Holand-Netherlands"
## [17] " Honduras" " Hong"
## [19] " Hungary" " India"
## [21] " Iran" " Ireland"
## [23] " Italy" " Jamaica"
## [25] " Japan" " Laos"
## [27] " Mexico" " Nicaragua"
## [29] " Outlying-US(Guam-USVI-etc)" " Peru"
## [31] " Philippines" " Poland"
## [33] " Portugal" " Puerto-Rico"
## [35] " Scotland" " South"
## [37] " Taiwan" " Thailand"
## [39] " Trinidad&Tobago" " United-States"
## [41] " Vietnam" " Yugoslavia"
##
## $income
## [1] " <=50K" " >50K"
```

```
# Convert " ?" data to NA and the remove rows with NA
income_data <- read.csv("adult.data",na.strings = c(" ?"),header = FALSE)

income_data <- na.omit(income_data)

# Set column names again
colnames(income_data) <- c("age","workclass","fnlwgt","education","education_num","marital_status",
                           "occupation","relationship","race","sex","capital_gain","capital_loss",
                           "hours_per_week","native_country","income")

# To get the dimension of new data
dim(income_data)
```

```
## [1] 30162 15
```

Splitting data into Train and validation sets:

We split the data into train and validation sets where validation set is 10% of the income_data. Train set has 27145 rows and validation has 3017 rows.

```

# Create Train and validation set with validation set having 10% data
set.seed(1, sample.kind="Rounding")
# if using R 3.5 or earlier, use `set.seed(1)` instead

test_index <- createDataPartition(income_data$income,p = 0.1,
                                  times = 1,list = FALSE)

trainset <- income_data[-test_index, ]
validation <- income_data[test_index,]

# To get dimension of train and validation data
dim(trainset)

## [1] 27145    15
dim(validation)

## [1] 3017     15

```

Data Exploration

There are 14 variables, so we explore each variable further.


1. age


The minimum and maximum age is 17 and 90 yrs respectively. Mean age is 38 yrs. Most data lies in the range 20 to 50 yrs. Age vs income plot shows that the proportion of income “>50K” is greater for age ≥ 30 and age ≤ 60 compared to other age groups.

```

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    17.00  28.00   37.00   38.43  47.00   90.00

```







2. Workclass

Private sector constitutes majority of workclass. The level “never-worked” has 0 entries whereas without pay is negligible. The same can be seen in histogram as well.

The barplot with workclass and income also shows that income “>50k” is maximum in private workclass.



##	Federal-gov	Local-gov	Never-worked	Private
##	855	1853	0	20028
##	Self-emp-inc	Self-emp-not-inc	State-gov	Without-pay
##	967	2268	1162	12



3. Final weight

Weights are assigned to participants based on demographic characteristics. For more details about `fnlwgt`, `adult.names` file can be referred. The minimum and maximum `fnlwgt` is 13769 and 1484705. As we can see in the boxplots there are many outliers and this variable is not very helpful in predicting income. Therefore we drop this variable.

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	13769	117526	178215	189453	237608	1484705





4. Education

The education of each individual can be seen in the histogram plot. Income “>50k” is more in the cases of bachelors, HS-grad, masters and some college. Doctorate and Prof-school have more cases of income “>50k”

compared to the total no. of cases in the respective classes.

##	10th	11th	12th	1st-4th	5th-6th
##	743	948	337	133	250
##	7th-8th	9th	Assoc-acdm	Assoc-voc	Bachelors
##	498	412	917	1158	4572
##	Doctorate	HS-grad	Masters	Preschool	Prof-school
##	328	8879	1481	42	481
##	Some-college				
##	5966				





5. Education Number

Education_num is the numerical form of education. Therefore this is a redundant variable and we drop this variable.

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1.00   9.00   10.00   10.13   13.00   16.00
```

```
## # A tibble: 16 x 3
## # Groups:   education_num [16]
##   education_num education      n
##         <int> <fct>    <int>
## 1             1 " Preschool"    42
## 2             2 " 1st-4th"   133
## 3             3 " 5th-6th"   250
## 4             4 " 7th-8th"   498
## 5             5 " 9th"     412
## 6             6 " 10th"    743
## 7             7 " 11th"    948
## 8             8 " 12th"    337
## 9             9 " HS-grad"  8879
## 10            10 " Some-college" 5966
## 11            11 " Assoc-voc"  1158
## 12            12 " Assoc-acdm"   917
## 13            13 " Bachelors"  4572
## 14            14 " Masters"   1481
## 15            15 " Prof-school"  481
## 16            16 " Doctorate"   328
```

6. Marital Status

Summary and barplots indicate that majority of individuals are of the category married-civ-spouse. This

same category has more no. of individuals with income ">50K".

6. Marital status


#Summary of marital_status

```
summary(trainset$marital_status)
```

```
##           Divorced      Married-AF-spouse      Married-civ-spouse
##           3772           19           12627
## Married-spouse-absent      Never-married           Separated
##           331           8795           851
##           Widowed
##           750
```


This plot displays marital_status distribution

```
trainset %>% ggplot(aes(marital_status)) +
  geom_bar(col = "black", fill = "grey") +
  theme(axis.text.x = element_text(angle=90,hjust = 1)) +
  ggtitle("Marital Status distribution")
```



This plot displays marital_status vs income



```
trainset %>% ggplot(aes(marital_status,fill = income)) +
  geom_bar() +
  theme(axis.text.x = element_text(angle=90,hjust = 1)) +
  ggtitle("Marital Status- Income distribution")
```



7. Occupation

People earning income “>50k” are mostly in the categories “Exec -Managerial” and “Prof-specialty”.

##	Adm-clerical	Armed-Forces	Craft-repair
##	3369	7	3620
##	Exec-managerial	Farming-fishing	Handlers-cleaners
##	3579	891	1208
##	Machine-op-inspct	Other-service	Priv-house-serv
##	1770	2910	121
##	Prof-specialty	Protective-serv	Sales
##	3649	577	3210
##	Tech-support	Transport-moving	
##	821	1413	


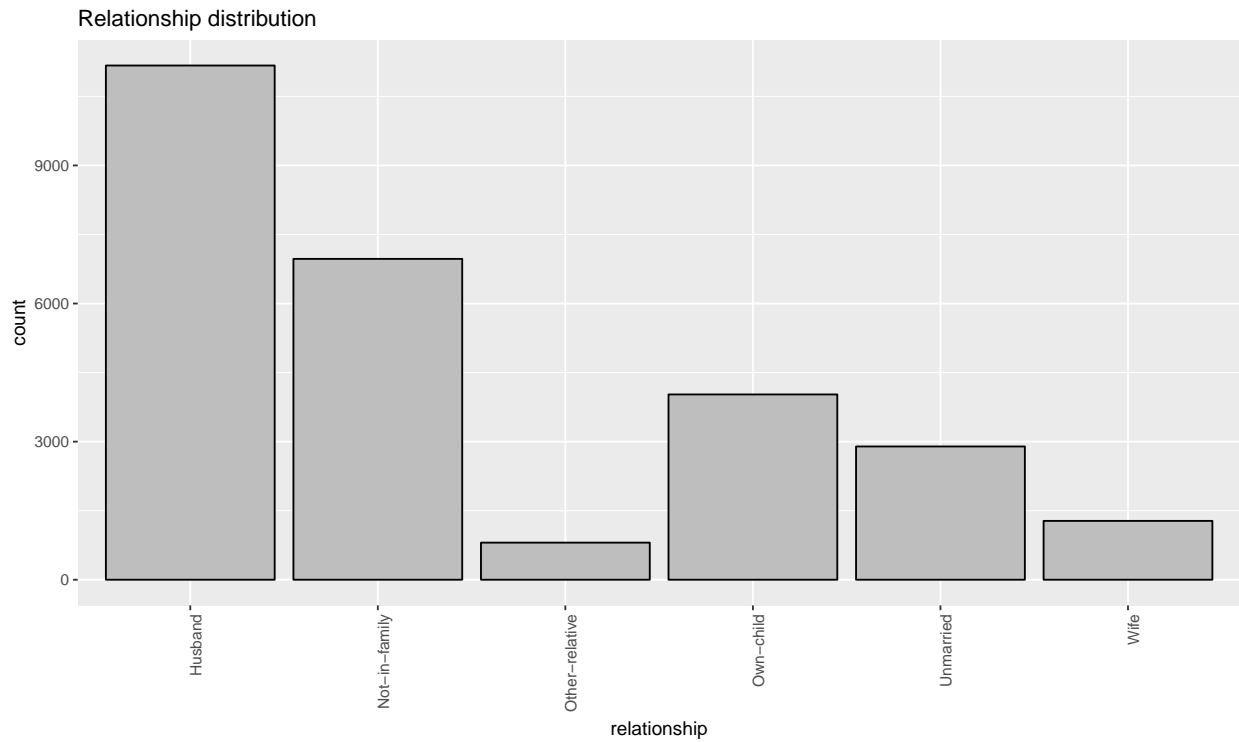


8. Relationship

“Husband” category has most no. of entries and also income “>50k” is mostly seen in husband category.

##	Husband	Not-in-family	Other-relative	Own-child
##	11171	6970	806	4025



```
##      Unmarried      Wife
##      2895      1278
```



9. Race

Most of the participants belong to “white” category followed by “black”.



##	Amer-Indian-Eskimo	Asian-Pac-Islander	Black
##	260	813	2537
##	Other	White	
##	206	23329	



10. Sex


Around 67% are males and 33% are females. No. of females with income “>50k” is less compared to males.

```
## Female    Male
##    8837   18308
```



Further we can observe that no. of males are more educated than females which explains the income



distribution difference in males and females.





11. Capital Gain and Capital Loss

Capital gain and Capital loss many values as “0” and there are many outliers as well. This is because not everybody invests in stock markets.

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##         0         0         0   1081     0   99999
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.00   0.00   0.00  88.89   0.00 4356.00
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

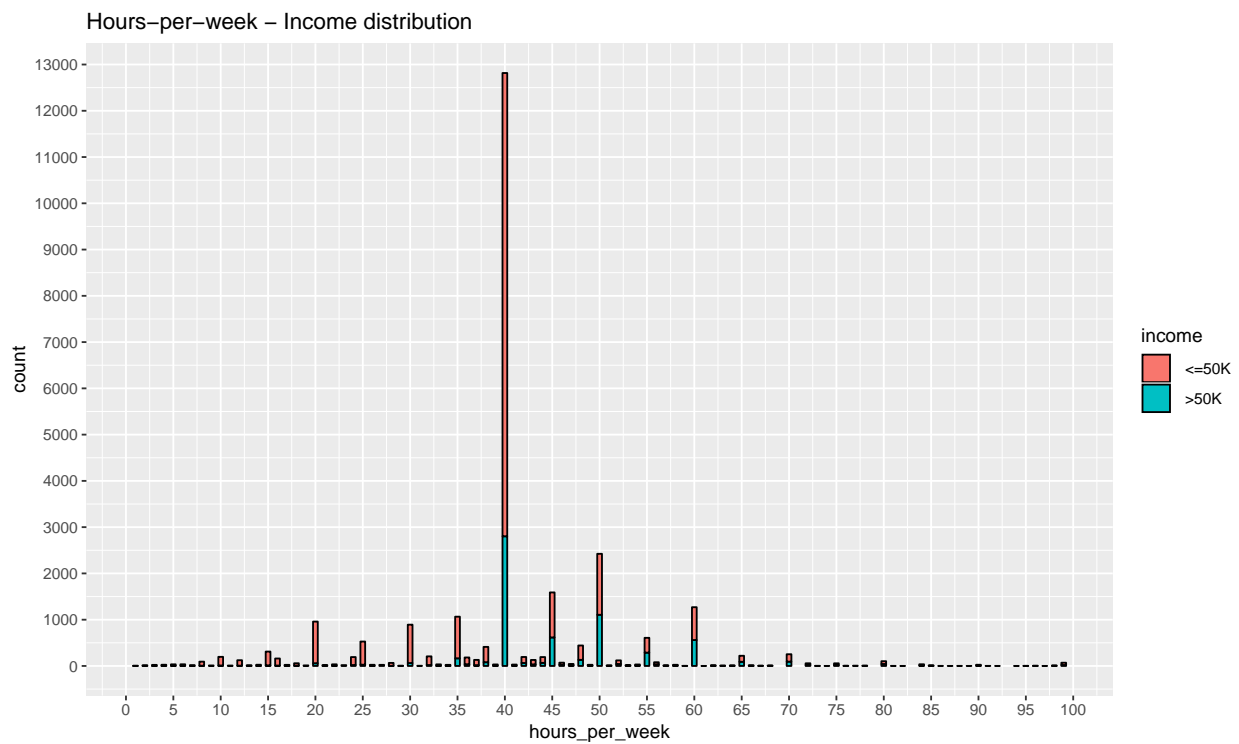
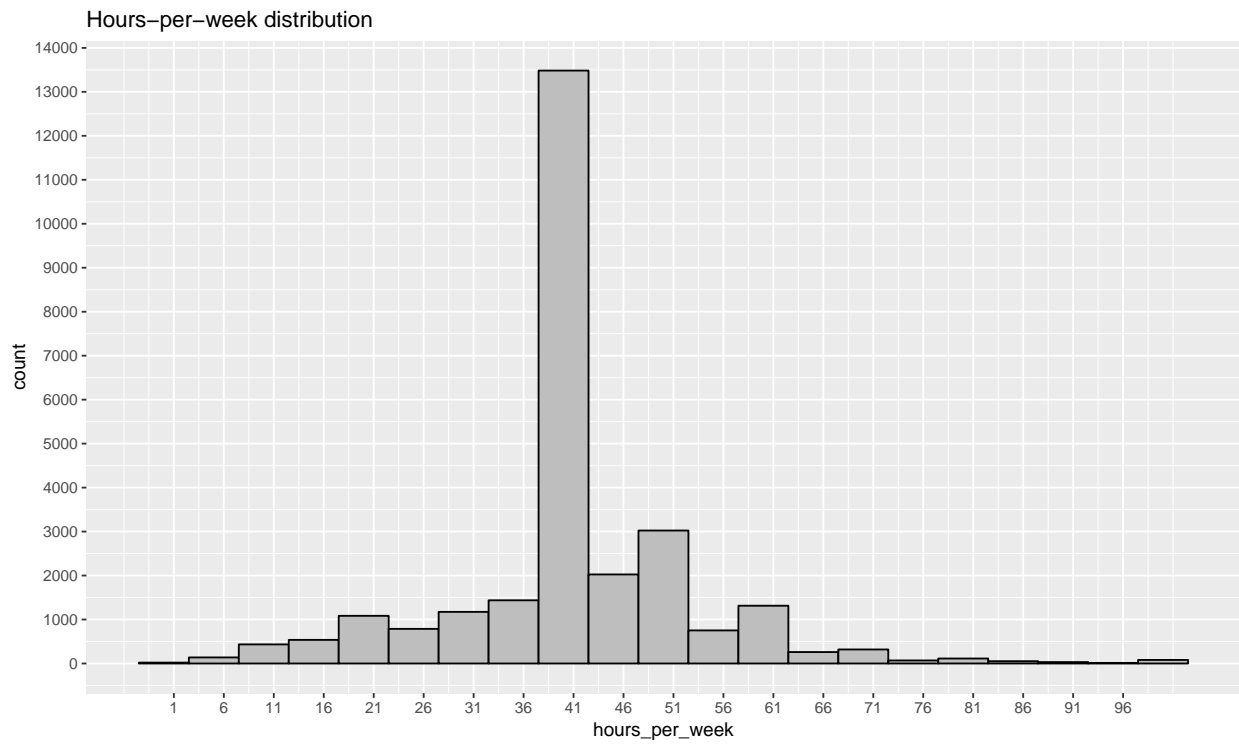


12. Hours-per-week

Minimum working hours is 1hr and maximum working hours is 99hrs. The mean no. of working hours is 40hrs. For income “>50k”, the mean working hours is between 45 to 50 hrs.

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
----	------	---------	--------	------	---------	------

1.00 40.00 40.00 40.92 45.00 99.00

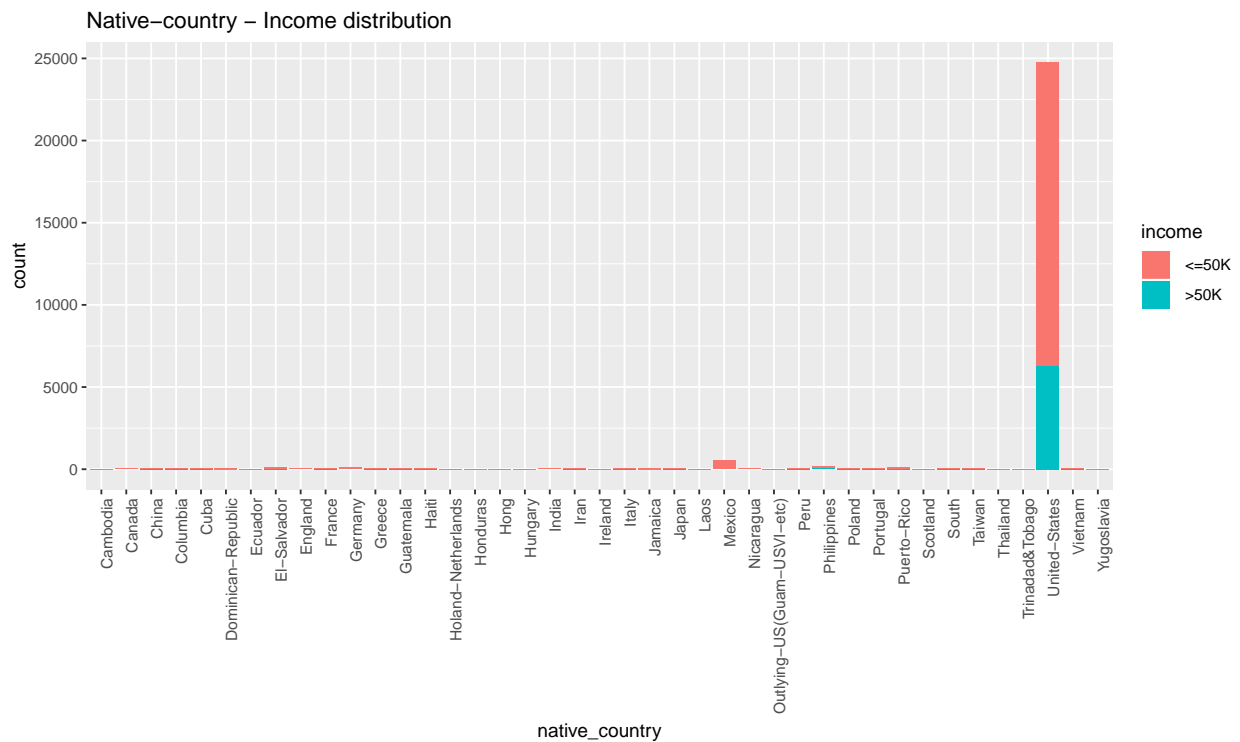



13. Native country

Summary shows that about 91% of the population belongs to United States.

Cambodia Canada

##	14	91
##	China	Columbia
##	59	52
##	Cuba	Dominican-Republic
##	84	63
##	Ecuador	El-Salvador
##	25	87
##	England	France
##	78	25
##	Germany	Greece
##	115	28
##	Guatemala	Haiti
##	56	41
##	Holand-Netherlands	Honduras
##	1	11
##	Hong	Hungary
##	18	12
##	India	Iran
##	81	40
##	Ireland	Italy
##	22	58
##	Jamaica	Japan
##	76	54
##	Laos	Mexico
##	16	543
##	Nicaragua	Outlying-US(Guam-USVI-etc)
##	31	14
##	Peru	Philippines
##	27	175
##	Poland	Portugal
##	51	32
##	Puerto-Rico	Scotland
##	100	8
##	South	Taiwan
##	67	38
##	Thailand	Trinidad&Tobago
##	17	16
##	United-States	Vietnam
##	24747	59
##	Yugoslavia	
##	13	



14. Income

This is the response variable. About 75% of the sample has income “<=50k”.

15. Income

```

# Summary of income
summary(trainset$income)

##      <=50K      >50K
##    20388     6757

# To calculate proportions for each factor level
percentage <- 100 * prop.table(table(trainset$income))

# Display in table format
cbind(freq = table(trainset$income), percentage = percentage)

##           freq percentage
## <=50K 20388     75.10775
## >50K  6757     24.89225

```

Data Modeling

We use all the variables except `fnlwgt` and `education_num` for data modeling. Algorithms used are as follows.

1. Logistic Regression

The response variable to be predicted is binary. So one of the best candidate algorithm is Logistic Regression.

Logistic Regression is suited for binary classification problems. Real-valued numbers are mapped to a value between 0 and 1 using Logistic Regression. Logistic Regression uses an equation where input values (x) are combined linearly using coefficient values (β) to predict an output value (y). Here, the output is in binary form (0 and 1). Each column used as input has an associated coefficient which is learned from training data.

We have used method as `glm` (generalized linear model) and family as binomial since this is Logistic Regression.

```

# 1. Logistic Regression

set.seed(1, sample.kind="Rounding")

# Fit data using caret package, method - glm, family - binomial
train_lr <- train(income ~ age + workclass + education + occupation +
                  relationship + hours_per_week + native_country +
                  race + sex + marital_status + capital_gain + capital_loss,
                  data=trainset,
                  method = "glm",
                  family="binomial")

# Predict income using the above fitted model
pred_lr <- predict(train_lr, validation)

# Save results of Confusion Matrix
lr_acc <- confusionMatrix(pred_lr, validation$income)

# Add results to a table
results <- tibble(Method="Logistic Regression",
                  Accuracy_Train = lr_acc$overall["Accuracy"],
                  F1_Train = lr_acc$byClass[7])
results %>% knitr::kable()

```


Method	Accuracy_Train	F1_Train
Logistic Regression	0.8578058	0.9081567

Accuracy obtained using logistic regression is 0.8578058 and F1 score is 0.9081567.

2. Support Vector Classification

Next we try with Support Vector Classification.

Support vector machine(SVM) algorithm finds a hyperplane in an N-dimensional space where N is the no. of features. This hyperplane distinctly classifies the data points. Our objective is to find a plane that has the maximum margin, i.e the maximum distance between the data points of both classes. The dimension of the hyperplane depends upon the number of features. Support vectors are data points that are closer to the hyperplane and influence the position and orientation of the hyperplane. Using these support vectors, we maximize the margin of the classifier. These are the points that help us build our SVM.

SVM from caret package takes a long time for computing, therefore we use svm from base package.

```
# 2. Support Vector Classifier

set.seed(3, sample.kind="Rounding")

## Fit data using svm from base package
svc <- svm(income ~ age + workclass + education + capital_gain +
           occupation + relationship + race + sex + capital_loss +
           hours_per_week + native_country + marital_status,
           data=trainset)

# Predict income using the above fitted model
pred_svc <- predict(svc, validation)

# Save results of Confusion Matrix
cm_svc <- confusionMatrix(pred_svc, validation$income)

# Add results to the results table
results <- bind_rows(results, tibble(Method="Support Vector Classifier",
                                     Accuracy_Train = cm_svc$overall["Accuracy"],
                                     F1_Train = cm_svc$byClass[7]))

results %>% knitr::kable()
```

Method	Accuracy_Train	F1_Train
Logistic Regression	0.8578058	0.9081567
Support Vector Classifier	0.8544912	0.9066950

Accuracy obtained using support vector classification is 0.8544912 and F1 score is 0.9066950.

3. Random Forest Classification

Random forest consists of a large number of individual decision trees. Random Forest is an ensemble model. Each individual tree in the random forest gives a class prediction and the class with the most votes becomes the model's prediction.

Random Forest from caret package takes a long time for computing, therefore we use rf from base package.

```

# 3. Random Forest Classifier

set.seed(4, sample.kind="Rounding")

## Fit data using rf from base package
raf <- randomForest(income ~ age + workclass + education + capital_gain +
                    occupation + relationship + race + sex + capital_loss +
                    hours_per_week + native_country + marital_status,
                    data = trainset)

# Predict income using the above fitted model
pred_raf <- predict(raf ,validation)

# Save results of Confusion Matrix
cm_raf <- confusionMatrix(pred_raf,validation$income)

# Add results to the results table
results <- bind_rows(results, tibble(Method="Random Forest Classifier",
                                     Accuracy_Train = cm_raf$overall["Accuracy"],
                                     F1_Train = cm_raf$byClass[7]))

results %>% knitr::kable()

```

Method	Accuracy_Train	F1_Train
Logistic Regression	0.8578058	0.9081567
Support Vector Classifier	0.8544912	0.9066950
Random Forest Classifier	0.8674180	0.9136069

Accuracy obtained using random forest classification is 0.8674180 and F1 score is 0.9136069. Here we see an improvement over logistic regression and svm.

4. Gradient Boosting Classifier

We use boosting method to convert weak learners into strong leaders. Here a new tree is fit based on the previous tree with modified weights. Gradient Boosting trains many models in a gradual, additive and sequential manner.

Accuracy obtained using gradient boosting classification is 0.8700696 and F1 score is 0.9169492. Here we see an improvement over logistic regression, svm and random forest classification.

Results

Now that we have trained our models, we apply them to test data. As noted earlier test data is available here: <https://archive.ics.uci.edu/ml/datasets/Adult> in adult.test file.

We follow the same preprocessing steps followed for train data set. The results are tabulated in the form of a table for easy comparison between models.

```

#-----
# Part 4: Results / Check the models on test set
#-----

# Downloading and Reading test data

```

```

if(!file.exists("adult.test")){
  download.file("http://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.test","adult.test")
}
income_test <- read.csv("adult.test",header = FALSE,skip=1)

# Set column names
colnames(income_test) <- c("age","workclass","fnlwgt","education","education_num","marital_status",
  "occupation","relationship","race","sex","capital_gain","capital_loss",
  "hours_per_week","native_country","income")

# explore test data

# To check the structure of test data
str(income_test)

## 'data.frame': 16281 obs. of 15 variables:
## $ age : int 25 38 28 44 18 34 29 63 24 55 ...
## $ workclass : Factor w/ 9 levels " ?"," Federal-gov",...: 5 5 3 5 1 5 1 7 5 5 ...
## $ fnlwgt : int 226802 89814 336951 160323 103497 198693 227026 104626 369667 104996 ...
## $ education : Factor w/ 16 levels " 10th"," 11th",...: 2 12 8 16 16 1 12 15 16 6 ...
## $ education_num : int 7 9 12 10 10 6 9 15 10 4 ...
## $ marital_status: Factor w/ 7 levels " Divorced"," Married-AF-spouse",...: 5 3 3 3 5 5 5 3 5 3 ...
## $ occupation : Factor w/ 15 levels " ?"," Adm-clerical",...: 8 6 12 8 1 9 1 11 9 4 ...
## $ relationship : Factor w/ 6 levels " Husband"," Not-in-family",...: 4 1 1 1 4 2 5 1 5 1 ...
## $ race : Factor w/ 5 levels " Amer-Indian-Eskimo",...: 3 5 5 3 5 5 3 5 5 5 ...
## $ sex : Factor w/ 2 levels " Female"," Male": 2 2 2 2 1 2 2 2 1 2 ...
## $ capital_gain : int 0 0 0 7688 0 0 0 3103 0 0 ...
## $ capital_loss : int 0 0 0 0 0 0 0 0 0 0 ...
## $ hours_per_week: int 40 50 40 40 30 30 40 32 40 10 ...
## $ native_country: Factor w/ 41 levels " ?"," Cambodia",...: 39 39 39 39 39 39 39 39 39 39 ...
## $ income : Factor w/ 2 levels " <=50K.", " >50K.": 1 1 2 2 1 1 1 2 1 1 ...

# To get the dimension of test data
dim(income_test)

## [1] 16281 15

# To get summary of test data
summary(income_test)

## age workclass fnlwgt
## Min. :17.00 Private :11210 Min. : 13492
## 1st Qu.:28.00 Self-emp-not-inc: 1321 1st Qu.: 116736
## Median :37.00 Local-gov : 1043 Median : 177831
## Mean :38.77 ? : 963 Mean : 189436
## 3rd Qu.:48.00 State-gov : 683 3rd Qu.: 238384
## Max. :90.00 Self-emp-inc : 579 Max. :1490400
## (Other) : 482
## education education_num marital_status
## HS-grad :5283 Min. : 1.00 Divorced :2190
## Some-college:3587 1st Qu.: 9.00 Married-AF-spouse : 14
## Bachelors :2670 Median :10.00 Married-civ-spouse :7403
## Masters : 934 Mean :10.07 Married-spouse-absent: 210
## Assoc-voc : 679 3rd Qu.:12.00 Never-married :5434
## 11th : 637 Max. :16.00 Separated : 505

```

```
## (Other) :2491 Widowed : 525
## occupation relationship
## Prof-specialty :2032 Husband :6523
## Exec-managerial:2020 Not-in-family :4278
## Craft-repair :2013 Other-relative: 525
## Sales :1854 Own-child :2513
## Adm-clerical :1841 Unmarried :1679
## Other-service :1628 Wife : 763
## (Other) :4893
## race sex capital_gain
## Amer-Indian-Eskimo: 159 Female: 5421 Min. : 0
## Asian-Pac-Islander: 480 Male :10860 1st Qu.: 0
## Black : 1561 Median : 0
## Other : 135 Mean : 1082
## White :13946 3rd Qu.: 0
## Max. :99999
##
## capital_loss hours_per_week native_country income
## Min. : 0.0 Min. : 1.00 United-States:14662 <=50K.:12435
## 1st Qu.: 0.0 1st Qu.:40.00 Mexico : 308 >50K. : 3846
## Median : 0.0 Median :40.00 ? : 274
## Mean : 87.9 Mean :40.39 Philippines : 97
## 3rd Qu.: 0.0 3rd Qu.:45.00 Puerto-Rico : 70
## Max. :3770.0 Max. :99.00 Germany : 69
## (Other) : 801
```

```
# To explore the levels in each column of test data
supply(income_test, levels)
```

```
## $age
## NULL
##
## $workclass
## [1] " ?" " Federal-gov" " Local-gov"
## [4] " Never-worked" " Private" " Self-emp-inc"
## [7] " Self-emp-not-inc" " State-gov" " Without-pay"
##
## $fnlwgt
## NULL
##
## $education
## [1] " 10th" " 11th" " 12th" " 1st-4th"
## [5] " 5th-6th" " 7th-8th" " 9th" " Assoc-acdm"
## [9] " Assoc-voc" " Bachelors" " Doctorate" " HS-grad"
## [13] " Masters" " Preschool" " Prof-school" " Some-college"
##
## $education_num
## NULL
##
## $marital_status
## [1] " Divorced" " Married-AF-spouse"
## [3] " Married-civ-spouse" " Married-spouse-absent"
## [5] " Never-married" " Separated"
## [7] " Widowed"
##
```

```

## $occupation
## [1] " ?" " Adm-clerical" " Armed-Forces"
## [4] " Craft-repair" " Exec-managerial" " Farming-fishing"
## [7] " Handlers-cleaners" " Machine-op-inspct" " Other-service"
## [10] " Priv-house-serv" " Prof-specialty" " Protective-serv"
## [13] " Sales" " Tech-support" " Transport-moving"
##
## $relationship
## [1] " Husband" " Not-in-family" " Other-relative" " Own-child"
## [5] " Unmarried" " Wife"
##
## $race
## [1] " Amer-Indian-Eskimo" " Asian-Pac-Islander" " Black"
## [4] " Other" " White"
##
## $sex
## [1] " Female" " Male"
##
## $capital_gain
## NULL
##
## $capital_loss
## NULL
##
## $hours_per_week
## NULL
##
## $native_country
## [1] " ?" " Cambodia"
## [3] " Canada" " China"
## [5] " Columbia" " Cuba"
## [7] " Dominican-Republic" " Ecuador"
## [9] " El-Salvador" " England"
## [11] " France" " Germany"
## [13] " Greece" " Guatemala"
## [15] " Haiti" " Honduras"
## [17] " Hong" " Hungary"
## [19] " India" " Iran"
## [21] " Ireland" " Italy"
## [23] " Jamaica" " Japan"
## [25] " Laos" " Mexico"
## [27] " Nicaragua" " Outlying-US(Guam-USVI-etc)"
## [29] " Peru" " Philippines"
## [31] " Poland" " Portugal"
## [33] " Puerto-Rico" " Scotland"
## [35] " South" " Taiwan"
## [37] " Thailand" " Trinidad&Tobago"
## [39] " United-States" " Vietnam"
## [41] " Yugoslavia"
##
## $income
## [1] " <=50K." " >50K."

```

```

# Convert " ?" data to NA and the remove rows with NA
income_test <- read.csv("adult.test",na.strings = c(" ?"),header = FALSE,skip=1)
income_test <- na.omit(income_test)

# Set column names again
colnames(income_test) <- c("age","workclass","fnlwgt","education","education_num","marital_status",
                           "occupation","relationship","race","sex","capital_gain","capital_loss",
                           "hours_per_week","native_country","income")

# Assign levels of income to test data
levels(income_test$income)[1] <- " <=50K"
levels(income_test$income)[2] <- " >50K"

# To ensure levels of native_country in train and test data are same
levels(income_test$native_country) <- levels(trainset$native_country)

# 1. Logistic Regression

# Use test data to predict income using the above fitted logistic regression model
pred_lrtest <- predict(train_lr,income_test)

# Save results of Confusion Matrix
lr_test <- confusionMatrix(pred_lrtest,income_test$income)

# Add results to the table
test_results <- tibble(Accuracy_Test = lr_test$overall["Accuracy"],
                       F1_Test = lr_test$byClass[7])
test_results %>% knitr::kable()

```

Accuracy_Test	F1_Test
0.8351262	0.8986985

```

# 2. Support Vector Classifier

set.seed(3, sample.kind="Rounding")

# Use test data to predict income using the above fitted svm model
pred_svctest <- predict(svc,income_test)

# Save results of Confusion Matrix
svc_test <- confusionMatrix(pred_svctest,income_test$income)

# Add results to the table
test_results <- bind_rows(test_results, tibble(
  Accuracy_Test = svc_test$overall["Accuracy"],
  F1_Test = svc_test$byClass[7]))
test_results %>% knitr::kable()

```

Accuracy_Test	F1_Test
0.8351262	0.8986985
0.8485392	0.9042039

3. Random Forest Classifier

```
set.seed(4, sample.kind="Rounding")

# Use test data to predict income using the above fitted random forest model
pred_raftest <- predict(raf,income_test)

# Save results of Confusion Matrix
raf_test <- confusionMatrix(pred_raftest,income_test$income)

# Add results to the table
test_results <- bind_rows(test_results, tibble(
  Accuracy_Test = raf_test$overall["Accuracy"],
  F1_Test = raf_test$byClass[7]))
test_results %>% knitr::kable()
```

Accuracy_Test	F1_Test
0.8351262	0.8986985
0.8485392	0.9042039
0.8324037	0.8852101

4. Gradient Boosting Classifier

```
set.seed(6, sample.kind="Rounding")

# Use test data to predict income using the above fitted gradient boosting model
pred_gbctest <- predict(gbc,income_test)

# Save results of Confusion Matrix
gbc_test <- confusionMatrix(pred_gbctest,income_test$income)

# Add results to the table
test_results <- bind_rows(test_results, tibble(
  Accuracy_Test = gbc_test$overall["Accuracy"],
  F1_Test= gbc_test$byClass[7]))
test_results %>% knitr::kable()
```

Accuracy_Test	F1_Test
0.8351262	0.8986985
0.8485392	0.9042039
0.8324037	0.8852101
0.8608234	0.9117399

```
# Add test results to train results table
results <- bind_cols(results,test_results)
results %>% knitr::kable()
```

Method	Accuracy_Train	F1_Train	Accuracy_Test	F1_Test
Logistic Regression	0.8578058	0.9081567	0.8351262	0.8986985
Support Vector Classifier	0.8544912	0.9066950	0.8485392	0.9042039

Method	Accuracy_Train	F1_Train	Accuracy_Test	F1_Test
Random Forest Classifier	0.8674180	0.9136069	0.8324037	0.8852101
Gradient Boosting Classifier	0.8700696	0.9169492	0.8608234	0.9117399

Out of the 4 models used on test data set, gradient boosting classifier has the best accuracy and F1 score followed by logistic regression. But gradient boosting classifier has more computation time compared to other models.

Conclusion

We use adult data set to predict whether income is “ $\leq 50k$ ” or “ $> 50k$ ”. After downloading data we removed all the rows with missing/unknown values. We explored data and realized final weight and education_num are not useful for the prediction. Therefore, we drop these variables.

Then we fit our data using the above models and on find that gradient boosing classifier has the best performance in this scenario but it also takes a longer time for the comptation.

The next comparable model is logistic regression which has better results along with lesser computaion time.

I would like to explore more on parameter tuning and also see if any processing can be done on the different levels of each variable. Maybe variables like capital gain and caplital loss can be fine tuned further and we can achieve even higher accuracy and F1 score.